

```

# (zero-shot) Instruction: Turn on a faucet by rotating its handle. The task
is finished when qpos of faucet handle is larger than target qpos.
import numpy as np
from scipy.spatial.distance import cdist

def compute_dense_reward(self, action):
    # Here I define the reward and penalty weights for different aspects of the task
    handle_reach_weight = 0.1
    grasp_handle_weight = 0.2
    rotation_weight = 0.7
    action_penalty_weight = 0.01

    # Here I define the grasp success reward
    grasp_success_reward = 0.1

    # Here I compute the distance between the robot's end effector and the faucet handle
    handle_pcd = self.faucet.handle.get_world_pcd()
    lfinger_cords = np.array([self.robot.lfinger.pose.p, ])
    rfinger_cords = np.array([self.robot.rfinger.pose.p, ])
    dist_lfinger_handle = cdist(lfinger_cords, handle_pcd).min(-1)[0]
    dist_rfinger_handle = cdist(rfinger_cords, handle_pcd).min(-1)[0]
    dist_handle_reach = max(dist_lfinger_handle, dist_rfinger_handle)
    handle_reach_reward = - handle_reach_weight * dist_handle_reach

    # Here I check if the robot has successfully grasped the faucet handle
    grasp_handle = self.robot.check_grasp(self.faucet.handle)
    grasp_handle_reward = grasp_handle_weight * grasp_handle if grasp_handle else 0

    # Here I calculate the rotation reward based on the difference between
    # the current and target joint position
    rotation_diff = max(0, self.faucet.handle.target_qpos - self.faucet.handle.qpos)
    rotation_reward = - rotation_weight * rotation_diff

    # Here I calculate the penalty for the robot's action
    action_penalty = - action_penalty_weight * np.linalg.norm(action)

    # The total reward is the sum of all the individual rewards and penalties
    reward = handle_reach_reward + grasp_handle_reward + rotation_reward + action_penalty

    if grasp_handle and rotation_diff == 0:
        reward += grasp_success_reward

    return reward

```

```

# (few-shot) Instruction: A dual-arm mobile robot needs to push a swivel chair
to a target location on the ground and prevent it from falling over.
import numpy as np
import scipy.spatial.distance as sdist

def compute_dense_reward(self, action):
    reward = -20.0

    actor = self.chair
    ee_coords = np.array(self.robot.get_ee_coords())
    ee_mids = np.array([ee_coords[:2].mean(0), ee_coords[2:].mean(0)])
    chair_pcd = self.chair.get_pcd()

    # EE approach chair
    dist_ees_to_chair = sdist.cdist(ee_coords, chair_pcd) # [4, N]
    dist_ees_to_chair = dist_ees_to_chair.min(1) # [4]
    dist_ee_to_chair = dist_ees_to_chair.mean()
    log_dist_ee_to_chair = np.log(dist_ee_to_chair + 1e-5)
    reward += - dist_ee_to_chair - np.clip(log_dist_ee_to_chair, -10, 0)

    # Penalize action
    action_norm = np.linalg.norm(action)
    reward -= action_norm * 1e-6

    # Keep chair standing
    z_axis_world = np.array([0, 0, 1])
    z_axis_chair = self.chair.pose.to_transformation_matrix()[[:3, 2]]
    chair_tilt = np.arccos(z_axis_chair[2])
    log_chair_tilt = np.log(chair_tilt + 1e-5)
    reward += -chair_tilt * 0.2

    # Chair velocity
    chair_vel = actor.velocity
    chair_vel_norm = np.linalg.norm(chair_vel)
    disp_chair_to_target = self.chair.pose.p[:2] - self.target_xy
    chair_vel_dir = sdist.cosine(chair_vel[:2], disp_chair_to_target)
    chair_ang_vel_norm = np.linalg.norm(actor.angular_velocity)

    # Stage reward
    stage_reward = 0

    dist_chair_to_target = np.linalg.norm(self.chair.pose.p[:2] - self.target_xy)

    if dist_ee_to_chair < 0.2:
        stage_reward += 2
    if dist_chair_to_target <= 0.3:
        stage_reward += 2
        reward += (np.exp(-chair_vel_norm * 10) * 2)
        if chair_vel_norm <= 0.1 and chair_ang_vel_norm <= 0.2:
            stage_reward += 2
            if chair_tilt <= 0.1 * np.pi:
                stage_reward += 2
    else:
        reward_vel = (chair_vel_dir - 1) * chair_vel_norm
        reward += np.clip(1 - np.exp(-reward_vel), -1, np.inf) * 2 - dist_chair_to_target * 2

    if chair_tilt > 0.4 * np.pi:
        stage_reward -= 2

    reward = reward + stage_reward
    return reward

```

```
# (zero-shot) Instruction: Pick up cube A and move it to the 3D goal position.
import numpy as np
```

```
def compute_dense_reward(self, action):
    # Normalize action
    action = np.clip(action, -1, 1)

    # Calculate distance between gripper and cube
    gripper_pos = self.robot.rr_pose.p
    cube_pos = self.cubeA.pose.p
    dist_gripper_cube = np.linalg.norm(gripper_pos - cube_pos)

    # Calculate distance between cube and goal
    goal_pos = self.goal_position
    dist_cube_goal = np.linalg.norm(goal_pos - cube_pos)

    # Check if the robot is grasping the cube
    grasping_cube = self.robot.check_grasp(self.cubeA)

    # Define reward components
    reward_dist_gripper_cube = -1.0 * dist_gripper_cube
    reward_dist_cube_goal = -1.0 * dist_cube_goal
    reward_grasping_cube = 1.0 if grasping_cube else -1.0

    # Define weights for reward components
    weight_dist_gripper_cube = 0.3
    weight_dist_cube_goal = 0.5
    weight_grasping_cube = 0.2

    # Calculate total reward
    reward = weight_dist_gripper_cube * reward_dist_gripper_cube \
        + weight_dist_cube_goal * reward_dist_cube_goal \
        + weight_grasping_cube * reward_grasping_cube

    # Regularization on action
    reward -= 0.01 * (action ** 2).sum()

    return reward
```

```
# (few-shot) Instruction: Pick up cube A and move it to the 3D goal position.
import numpy as np
```

```
def compute_dense_reward(self, action):
    reward = 0.0

    # Check if the cube is placed at the goal position and the robot is static
    cube_at_goal = np.linalg.norm(self.cubeA.pose.p - self.goal_position) \
        <= self.cube_half_size
    is_robot_static = np.max(np.abs(self.robot.qvel)) <= 0.2

    # If the cube is placed at the goal and the robot is static, return a high reward
    if cube_at_goal and is_robot_static:
        reward += 2.25
        return reward

    # reaching reward, encourages the robot to reach the cube
    gripper_pos = self.robot.rr_pose.p
    obj_pos = self.cubeA.pose.p
    dist_to_obj = np.linalg.norm(gripper_pos - obj_pos)
    reaching_reward = 1 - np.tanh(5 * dist_to_obj)
    reward += reaching_reward

    # grasp reward, encourages the robot to grasp the cube
    is_grasped = self.robot.check_grasp(self.cubeA, max_angle=30)
    if is_grasped:
        reward += 0.25

    # placement reward, encourages the robot to place the cube at the goal
    if is_grasped:
        dist_to_goal = np.linalg.norm(self.cubeA.pose.p - self.goal_position)
        placement_reward = 1 - np.tanh(5 * dist_to_goal)
        reward += placement_reward

    # regularization term on robot's action
    action_reg = -np.sum(np.square(action)) / len(action)
    reward += 0.1 * action_reg

    return reward
```

(few-shot interactive) Instruction: Pick up cubeA and place it on cubeB. The task is finished when cubeA is on top of cubeB stably (i.e. cubeA is static) and isn't grasped by the gripper.

```
import numpy as np

def compute_dense_reward(self, action):
    reward = 0.0
    # Check if cube A is on top of cube B and whether it is stable
    is_obj_on_target = np.linalg.norm(self.cubeA.pose.p - (self.cubeB.pose.p \
        + np.array([0, 0, self.cube_half_size*2]))) <= 0.025
    is_obj_static = self.cubeA.check_static()
    is_grasped = self.robot.check_grasp(self.cubeA)
    + gripper_openness = self.robot.gripper_openness

- success = is_obj_on_target and is_obj_static and not is_grasped
+ success = is_obj_on_target and is_obj_static

    if success:
        reward += 5
- return reward
+ if not is_grasped:
+     reward += 
+     return reward

    # Compute distance from the robot's end-effector to cube A
    tcp_to_obj_pos = self.cubeA.pose.p - self.robot.ee_pose.p
    tcp_to_obj_dist = np.linalg.norm(tcp_to_obj_pos)
    reaching_reward = 1 - np.tanh(5 * tcp_to_obj_dist)
    reward += reaching_reward

    # Check if the robot has grasped cube A
    reward += 1 if is_grasped else 0.0
    if is_grasped:
        # Compute distance from cube A to the top of cube B
        obj_to_goal_pos = self.cubeB.pose.p + np.array([0, 0, self.cube_half_size*2]) \
            - self.cubeA.pose.p
        obj_to_goal_dist = np.linalg.norm(obj_to_goal_pos)
        place_reward = 1 - np.tanh(5 * obj_to_goal_dist)
        reward += place_reward

+ if is_obj_on_target and is_obj_static:
+     # Encourage robot to release grasp when cube A is on top of cube B and stable
+     release_reward = 3 if not is_grasped and gripper_openness > 0.5 else -
+     reward += release_reward
+ else:
+     # Encourage the robot to keep the grasp if it's not on target
+     holding_reward = 1 if is_grasped else -1
+     reward += holding_reward
+ else:
+     if not is_obj_on_target:
+         # Encourage the robot to move closer to the target when not grasping
+         tcp_to_target_pos = self.cubeB.pose.p + np.array([0, 0, self.cube_half_size*2]) \
+             - self.robot.ee_pose.p
+         tcp_to_target_dist = np.linalg.norm(tcp_to_target_pos)
+         target_reward = 1 - np.tanh(5 * tcp_to_target_dist)
+         reward += target_reward

    return reward
```