

```

# (zero-shot) Instruction: Turn on a faucet by rotating its handle. The task
is finished when qpos of faucet handle is larger than target qpos.
import numpy as np
from scipy.spatial.distance import cdist

def compute_dense_reward(self, action):
    # Here I define the reward and penalty weights for different aspects of the task
    handle_reach_weight = 0.1
    grasp_handle_weight = 0.2
    rotation_weight = 0.7
    action_penalty_weight = 0.01

    # Here I define the grasp success reward
    grasp_success_reward = 0.1

    # Here I compute the distance between the robot's end effector and the faucet handle
    handle_pcd = self.faucet.handle.get_world_pcd()
    lfinger_cords = np.array([self.robot.lfinger.pose.p, ])
    rfinger_cords = np.array([self.robot.rfinger.pose.p, ])
    dist_lfinger_handle = cdist(lfinger_cords, handle_pcd).min(-1)[0]
    dist_rfinger_handle = cdist(rfinger_cords, handle_pcd).min(-1)[0]
    dist_handle_reach = max(dist_lfinger_handle, dist_rfinger_handle)
    handle_reach_reward = - handle_reach_weight * dist_handle_reach

    # Here I check if the robot has successfully grasped the faucet handle
    grasp_handle = self.robot.check_grasp(self.faucet.handle)
    grasp_handle_reward = grasp_handle_weight * grasp_handle if grasp_handle else 0

    # Here I calculate the rotation reward based on the difference between
    # the current and target joint position
    rotation_diff = max(0, self.faucet.handle.target_qpos - self.faucet.handle.qpos)
    rotation_reward = - rotation_weight * rotation_diff

    # Here I calculate the penalty for the robot's action
    action_penalty = - action_penalty_weight * np.linalg.norm(action)

    # The total reward is the sum of all the individual rewards and penalties
    reward = handle_reach_reward + grasp_handle_reward + rotation_reward + action_penalty

    if grasp_handle and rotation_diff == 0:
        reward += grasp_success_reward

    return reward

```

```

# (few-shot) Instruction: A dual-arm mobile robot needs to push a swivel chair
to a target location on the ground and prevent it from falling over.
import numpy as np
import scipy.spatial.distance as sdist

def compute_dense_reward(self, action):
    reward = -20.0

    actor = self.chair
    ee_coords = np.array(self.robot.get_ee_coords())
    ee_mids = np.array([ee_coords[2].mean(0), ee_coords[2].mean(0)])
    chair_pcd = self.chair.get_pcd()

    # EE approach chair
    dist_ees_to_chair = sdist.cdist(ee_coords, chair_pcd) # [4, N]
    dist_ees_to_chair = dist_ees_to_chair.min(1) # [4]
    dist_ee_to_chair = dist_ees_to_chair.mean()
    log_dist_ee_to_chair = np.log(dist_ee_to_chair + 1e-5)
    reward += - dist_ee_to_chair - np.clip(log_dist_ee_to_chair, -10, 0)

    # Penalize action
    action_norm = np.linalg.norm(action)
    reward -= action_norm * 1e-6

    # Keep chair standing
    z_axis_world = np.array([0, 0, 1])
    z_axis_chair = self.chair.pose.to_transformation_matrix()[3, 2]
    chair_tilt = np.arccos(z_axis_chair[2])
    log_chair_tilt = np.log(chair_tilt + 1e-5)
    reward += -chair_tilt * 0.2

    # Chair velocity
    chair_vel = actor.velocity
    chair_vel_norm = np.linalg.norm(chair_vel)
    disp_chair_to_target = self.chair.pose.p[2] - self.target_xy
    chair_vel_dir = sdist.cosine(chair_vel[2], disp_chair_to_target)
    chair_ang_vel_norm = np.linalg.norm(actor.angular_velocity)

    # Stage reward
    stage_reward = 0

    dist_chair_to_target = np.linalg.norm(self.chair.pose.p[2]-self.target_xy)

    if dist_ee_to_chair < 0.2:
        stage_reward += 2
    if dist_chair_to_target <= 0.3:
        stage_reward += 2
        reward += (np.exp(-chair_vel_norm * 10) * 2)
        if chair_vel_norm <= 0.1 and chair_ang_vel_norm <= 0.2:
            stage_reward += 2
            if chair_tilt <= 0.1 * np.pi:
                stage_reward += 2
    else:
        reward_vel = (chair_vel_dir - 1) * chair_vel_norm
        reward += np.clip(1-np.exp(-reward_vel), -1, np.inf)*2-dist_chair_to_target*2

    if chair_tilt > 0.4 * np.pi:
        stage_reward -= 2

    reward = reward + stage_reward
    return reward

```

```
# (zero-shot) Instruction: Pick up cube A and move it to the 3D goal position.
import numpy as np
```

```
def compute_dense_reward(self, action):
    # Normalize action
    action = np.clip(action, -1, 1)

    # Calculate distance between gripper and cube
    gripper_pos = self.robot.rr_pose.p
    cube_pos = self.cubeA.pose.p
    dist_gripper_cube = np.linalg.norm(gripper_pos - cube_pos)

    # Calculate distance between cube and goal
    goal_pos = self.goal_position
    dist_cube_goal = np.linalg.norm(goal_pos - cube_pos)

    # Check if the robot is grasping the cube
    grasping_cube = self.robot.check_grasp(self.cubeA)

    # Define reward components
    reward_dist_gripper_cube = -1.0 * dist_gripper_cube
    reward_dist_cube_goal = -1.0 * dist_cube_goal
    reward_grasping_cube = 1.0 if grasping_cube else -1.0

    # Define weights for reward components
    weight_dist_gripper_cube = 0.3
    weight_dist_cube_goal = 0.5
    weight_grasping_cube = 0.2

    # Calculate total reward
    reward = weight_dist_gripper_cube * reward_dist_gripper_cube \
        + weight_dist_cube_goal * reward_dist_cube_goal \
        + weight_grasping_cube * reward_grasping_cube

    # Regularization on action
    reward -= 0.01 * (action ** 2).sum()

    return reward
```

```
# (few-shot) Instruction: Pick up cube A and move it to the 3D goal position.
import numpy as np
```

```
def compute_dense_reward(self, action):
    reward = 0.0

    # Check if the cube is placed at the goal position and the robot is static
    cube_at_goal = np.linalg.norm(self.cubeA.pose.p - self.goal_position) \
        <= self.cube_half_size
    is_robot_static = np.max(np.abs(self.robot.qvel)) <= 0.2

    # If the cube is placed at the goal and the robot is static, return a high reward
    if cube_at_goal and is_robot_static:
        reward += 2.25
        return reward

    # reaching reward, encourages the robot to reach the cube
    gripper_pos = self.robot.rr_pose.p
    obj_pos = self.cubeA.pose.p
    dist_to_obj = np.linalg.norm(gripper_pos - obj_pos)
    reaching_reward = 1 - np.tanh(5 * dist_to_obj)
    reward += reaching_reward

    # grasp reward, encourages the robot to grasp the cube
    is_grasped = self.robot.check_grasp(self.cubeA, max_angle=30)
    if is_grasped:
        reward += 0.25

    # placement reward, encourages the robot to place the cube at the goal
    if is_grasped:
        dist_to_goal = np.linalg.norm(self.cubeA.pose.p - self.goal_position)
        placement_reward = 1 - np.tanh(5 * dist_to_goal)
        reward += placement_reward

    # regularization term on robot's action
    action_reg = -np.sum(np.square(action)) / len(action)
    reward += 0.1 * action_reg

    return reward
```